

CS2011 (Intro To Computer Systems)

Lab4: Cache Simulator

1 Logistics

This is an individual project. You must run this lab on a 64-bit x86-64 machine. You are encouraged to use the same VM you were asked to install in previous labs (i.e., Ubuntu 20.04 or Linux Mint 20). Remember that your submission will only be tested on an Ubuntu 20.04 or Linux Mint 20 OS (Operating System) with compiler configurations and libraries similar to those that come preinstalled or installed later via the default repository on those OSes. Thus we advise you to implement and test your code on such machines/VMs before submitting it.

2 Overview

This lab will help you understand the impact that cache memories can have on the performance of your C programs. You will write a small C program (about 200-300 lines) that simulates the behavior of a cache memory.

3 Downloading the assignment

Your lab materials are contained in a tar archive file called `lab4-handout.tar`, which you can download from Canvas.

Start by copying `lab4-handout.tar` to a directory on your VM in which you plan to do your work. Then open your terminal and navigate to that directory using the `'cd'` command. Then run the following command to extract the archive:

```
linux> tar xvf lab4-handout.tar
```

This will create a directory called `lab4-handout` that contains a number of files. You will be modifying the `csm.c` file only. To compile this file, type:

```
linux> make clean
linux> make
```

WARNING: Do not let the Windows WinZip program open up your `.tar` file (many Web browsers are set to do this automatically). Instead, save the file to your Linux directory and use the Linux `tar` program to extract the files. In general, for this class you should NEVER use any platform other than Linux to modify your files. Doing so can cause loss of data (and important work!).

4 Description

In this lab, you will implement a cache simulator.

4.1 Reference Trace Files

The `traces` subdirectory of the handout directory contains a collection of *reference trace files* that we will use to evaluate the correctness of the cache simulator you will write. The trace files are generated by a Linux program called `valgrind`. To install `valgrind`, run the following command in your terminal:

```
linux> sudo apt upgrade
linux> sudo apt install valgrind
```

Then, for example, typing

```
linux> valgrind --log-fd=1 --tool=lackey -v --trace-mem=yes ls -l
```

on the command line runs the executable program “`ls -l`”, captures a trace of each of its memory accesses in the order they occur, and prints them on `stdout`.

Valgrind memory traces have the following form:

```
I 0400d7d4,8
M 0421c7f0,4
L 04f6b868,8
S 7ff0005c8,8
```

Each line denotes one or two memory accesses. The format of each line is

```
[space]operation address,size
```

The *operation* field denotes the type of memory access:

“I” denotes an instruction load,

“L” a data load,

“S” a data store, and

“M” a data modify (i.e., a data load followed by a data store).

There is never a space before each “I”. There is always a space before each “M”, “L”, and “S”.

The *address* field specifies a 64-bit hexadecimal memory address.

The *size* field specifies the number of bytes accessed by the operation.

4.2 Writing the Cache Simulator

You are asked to write a cache simulator in `csim.c` that takes a `valgrind` memory trace as input, simulates the hit/miss behavior of a cache memory on this trace, and outputs the total number of hits, misses, and evictions.

We have provided you with the binary executable of a *reference cache simulator*, called `csim-ref`, that simulates the behavior of a cache with arbitrary size and associativity on a `valgrind` trace file. It uses the LRU (least-recently used) replacement policy when choosing which cache line to evict.

The reference simulator takes the following command-line arguments:

Usage: `./csim-ref [-hv] -s <s> -E <E> -b -t <tracefile>`

- `-h`: Optional help flag that prints usage info
- `-v`: Optional verbose flag that displays trace info
- `-s <s>`: Number of set index bits ($S = 2^s$ is the number of sets)
- `-E <E>`: Associativity (number of lines per set)
- `-b `: Number of block bits ($B = 2^b$ is the block size)
- `-t <tracefile>`: Name of the `valgrind` trace to replay

The command-line arguments are based on the notation (s , E , and b) from the textbook. For example:

```
linux> ./csim-ref -s 4 -E 1 -b 4 -t traces/yi.trace
hits:4 misses:5 evictions:3
```

The same example in verbose mode:

```
linux> ./csim-ref -v -s 4 -E 1 -b 4 -t traces/yi.trace
L 10,1 miss
M 20,1 miss hit
L 22,1 hit
S 18,1 hit
L 110,1 miss eviction
L 210,1 miss eviction
M 12,1 miss eviction hit
hits:4 misses:5 evictions:3
```

Your job is to fill in the `csim.c` file so that it takes the same command line arguments and produces the identical output as the reference simulator. Notice that this file is almost completely empty. You'll need to write it from scratch.

Programming Rules

- Include your name and UC ID in the header comment for `csim.c`.
- Your `csim.c` file must compile without warnings in order to receive credit.
- Your simulator must work correctly for arbitrary s , E , and b . This means that you will need to allocate storage for your simulator's data structures using the `malloc` function. Type “man malloc” for information about this function.
- For this lab, we are interested only in data cache performance, so your simulator should ignore all instruction cache accesses (lines starting with “I”). Recall that `valgrind` always puts “I” in the first column (with no preceding space), and “M”, “L”, and “S” in the second column (with a preceding space). This may help you parse the trace.
- To receive credit, you must call the function `printSummary`, with the total number of hits, misses, and evictions, at the end of your `main` function:

```
printSummary(hit_count, miss_count, eviction_count);
```

- For this lab, you should assume that memory accesses are aligned properly, such that a single memory access never crosses block boundaries. By making this assumption, you can ignore the request sizes in the `valgrind` traces.

5 Evaluation

This section describes how your work will be evaluated. The full score for this lab is 35 points:

- Code: 27 Points
- Style: 8 Points

5.1 Code Evaluation

We will run your cache simulator using different cache parameters and traces. There are eight test cases, each worth 3 points, except for the last case, which is worth 6 points:

```
linux> ./csim -s 1 -E 1 -b 1 -t traces/yi2.trace
linux> ./csim -s 4 -E 2 -b 4 -t traces/yi.trace
linux> ./csim -s 2 -E 1 -b 4 -t traces/dave.trace
linux> ./csim -s 2 -E 1 -b 3 -t traces/trans.trace
linux> ./csim -s 2 -E 2 -b 3 -t traces/trans.trace
linux> ./csim -s 2 -E 4 -b 3 -t traces/trans.trace
linux> ./csim -s 5 -E 1 -b 5 -t traces/trans.trace
linux> ./csim -s 5 -E 1 -b 5 -t traces/long.trace
```

You can use the reference simulator `csim-ref` to obtain the correct answer for each of these test cases. During debugging, use the `-v` option for a detailed record of each hit and miss.

For each test case, outputting the correct number of cache hits, misses and evictions will give you full credit for that test case. Each of your reported number of hits, misses and evictions is worth 1/3 of the credit for that test case. That is, if a particular test case is worth 3 points, and your simulator outputs the correct number of hits and misses, but reports the wrong number of evictions, then you will earn 2 points.

5.2 Evaluation for Style

There are 8 points for coding style. These will be assigned manually by your grader.

The style and format of your code will be taken seriously when grading. The grader will inspect your code for excessive usage of local variables, lack of comments (comments for each function, comments for important lines of code, etc), the lack of code to free memory after use, the lack of code to handle all command line options passed to your `csim` program, etc.

6 Working on the Lab

You are provided with an autograding program, called `test-csim`, that tests the correctness of your cache simulator on the reference traces. Be sure to compile your simulator before running the test:

```
linux> make
linux> ./test-csim
```

Points	(s,E,b)	Your simulator			Reference simulator			
		Hits	Misses	Evicts	Hits	Misses	Evicts	
3	(1,1,1)	9	8	6	9	8	6	traces/yi2.trace
3	(4,2,4)	4	5	2	4	5	2	traces/yi.trace
3	(2,1,4)	2	3	1	2	3	1	traces/dave.trace
3	(2,1,3)	167	71	67	167	71	67	traces/trans.trace
3	(2,2,3)	201	37	29	201	37	29	traces/trans.trace
3	(2,4,3)	212	26	10	212	26	10	traces/trans.trace
3	(5,1,5)	231	7	0	231	7	0	traces/trans.trace
6	(5,1,5)	265189	21775	21743	265189	21775	21743	traces/long.trace

27

For each test, it shows the number of points you earned, the cache parameters, the input trace file, and a comparison of the results from your simulator and the reference simulator.

Here are some hints and suggestions while working on the lab:

- Do your initial debugging on the small traces, such as `traces/dave.trace`.
- The reference simulator takes an optional `-v` argument that enables verbose output, displaying the hits, misses, and evictions that occur as a result of each memory access. You are not required to implement this feature in your `csim.c` code, but we strongly recommend that you do so. It will

help you debug by allowing you to directly compare the behavior of your simulator with the reference simulator on the reference trace files.

- We recommend that you use the `getopt` function to parse your command line arguments. You'll need the following header files:

```
#include <getopt.h>
#include <stdlib.h>
#include <unistd.h>
```

See “man 3 getopt” for details.

- Each data load (L) or store (S) operation can cause at most one cache miss. The data modify operation (M) is treated as a load followed by a store to the same address. Thus, an M operation can result in two cache hits, or a miss and a hit plus a possible eviction.

6.1 Driver Program

We have provided you with a *driver program*, called `driver.py`, that performs a complete evaluation of your simulator. This is the same program your grader uses to evaluate your handins. The driver uses `test-csim` to evaluate your simulator. Then it prints a summary of your results and the points you have earned.

To run the driver, type:

```
linux> python driver.py
```

You need to have python 2.7 installed. If it is not, run the following commands to install it:

```
linux> sudo apt upgrade
linux> sudo apt install python
```

7 Handing in Your Work

To receive a score, you should upload your submission to Canvas. You may handin as often as you like until the due date but please keep the number of submissions as MINIMAL as possible. The last submission is what will be graded. Submit a tar file containing your `csim.c` implementation to Canvas. Running the `make` command from the command line will generate the tar file, `lab4-handin.tar`, for you. You can upload this file to Canvas. IMPORTANT: Do not assume your submission will succeed! You should ALWAYS check to see if your submission was successful on Canvas. IMPORTANT: Do not upload files in other archive formats, such as those with extensions `.zip`, `.gzip`, or `.tgz`.